

```
In [ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix

from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

url = "/content/drive/MyDrive/WA_Fn-UseC_-Telco-Customer-Churn.csv"
data = pd.read_csv(url)
```

```
In [ ]: def plot_cf_matrix(y_test, y_pred, target_names):
    cf_matrix = confusion_matrix(y_true=y_test, y_pred=y_pred)
    df_matrix_df = pd.DataFrame(cf_matrix, index = target_names, columns = target_names)
    sns.heatmap(df_matrix_df, annot=True, cmap='RdYlBu', linewidth=.5, fmt='d')
    plt.title('Confusion matrix', y=1.1)
    plt.ylabel('Actual label')
    plt.xlabel('Predicted label')

target_names = ['No Churn', 'Churn']
```

```
In [ ]: print("特徴類型表格 : ")
print(data.dtypes)
```

特徵類型表格：

customerID	object
gender	object
SeniorCitizen	int64
Partner	object
Dependents	object
tenure	int64
PhoneService	object
MultipleLines	object
InternetService	object
OnlineSecurity	object
OnlineBackup	object
DeviceProtection	object
TechSupport	object
StreamingTV	object
StreamingMovies	object
Contract	object
PaperlessBilling	object
PaymentMethod	object
MonthlyCharges	float64
TotalCharges	object
Churn	object
dtype:	object

```
In [ ]: print("數據的形狀 ( 行數和列數 ) :", data.shape)
```

```
print("\n前五行數據 : ")  
print(data.head())
```

數據的形狀 ( 行數和列數 ) : (7043, 21)

前五行數據 :

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService
0	7590-VHVEG	Female	0	Yes	No	1	No
1	5575-GNVDE	Male	0	No	No	34	Yes
2	3668-QPYBK	Male	0	No	No	2	Yes
3	7795-CFOCW	Male	0	No	No	45	No
4	9237-HQITU	Female	0	No	No	2	Yes

	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection
0	No phone service	DSL	No	...	No
1	No	DSL	Yes	...	Yes
2	No	DSL	Yes	...	No
3	No phone service	DSL	Yes	...	Yes
4	No	Fiber optic	No	...	No


	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling
0	No	No	No	Month-to-month	Yes
1	No	No	No	One year	No
2	No	No	No	Month-to-month	Yes
3	Yes	No	No	One year	No
4	No	No	No	Month-to-month	Yes

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

```
In [ ]: # 繪製Churn比例的圓餅圖
churn_counts = data["Churn"].value_counts()
plt.figure(figsize=(6, 6))
plt.pie(churn_counts, labels=churn_counts.index, autopct="%1.1f%%")
plt.title("Churn Rate")
plt.show()
```


## Churn Rate



```
In [ ]: # 繪製特徵之間的相關性熱力圖
plt.figure(figsize=(16, 8))
sns.heatmap(data.corr(), annot=True, cmap="coolwarm", linewidths=0.5, fmt='
plt.title("Correlation Heatmap")
plt.show()
```

<ipython-input-5-d76f5fd7cc90>:3: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
sns.heatmap(data.corr(), annot=True, cmap="coolwarm", linewidths=0.5, fmt=".2f")
```



```
In [ ]: for column in data.columns:
        if data[column].dtype == object:
            encoder = LabelEncoder()
            data[column] = encoder.fit_transform(data[column])
```

```
In [ ]: X = data.drop(["Churn", 'customerID'], axis=1)
        y = data["Churn"]

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
```

```
In [ ]: from sklearn.feature_selection import chi2
```

```
In [ ]: chi2, p_value = chi2(X_train, y_train)
```

```
In [ ]: chi_df = pd.DataFrame({ 'chi-squares test': chi2, 'p-value':p_value})
        chi_df['p-value'] = chi_df['p-value'].round(4)
        chi_df = chi_df.sort_values('chi-squares test', ascending=False)
        chi_df.style.background_gradient(cmap='Blues')
```

Out[ ]:

	chi-squares test	p-value
4	12349.737573	0.000000
17	2795.912582	0.000000
18	2234.827250	0.000000
14	873.211647	0.000000
8	439.328373	0.000000
11	407.908476	0.000000
9	177.543358	0.000000
10	144.973153	0.000000
3	104.146764	0.000000
1	102.275240	0.000000
15	84.520367	0.000000
2	58.775844	0.000000
16	41.581329	0.000000
7	7.813628	0.005200
6	6.523859	0.010600
12	6.426778	0.011200
13	4.752568	0.029300
0	0.377925	0.538700
5	0.043591	0.834600

```
In [ ]: from sklearn.feature_selection import SelectKBest, chi2
```

```
chi2Filter = SelectKBest(chi2, k=8)
chi2Filter.fit(X_train, y_train)

# 選取完的 top 20 特徵
X_train_new = chi2Filter.transform(X_train)
X_test_new = chi2Filter.transform(X_test)
```

```
In [ ]: chi2Filter.get_feature_names_out()
```

```
Out[ ]: array(['tenure', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
               'TechSupport', 'Contract', 'MonthlyCharges', 'TotalCharges'],
              dtype=object)
```

```
In [ ]: from sklearn.svm import SVC
```

```

rf = RandomForestClassifier(random_state=42)

rf.fit(X_train_new, y_train)

chi2_select_pred_20 = rf.predict(X_test_new)

report = classification_report(y_test, chi2_select_pred_20)
print(report)

```

	precision	recall	f1-score	support
0	0.83	0.90	0.86	1036
1	0.64	0.49	0.55	373
accuracy			0.79	1409
macro avg	0.73	0.70	0.71	1409
weighted avg	0.78	0.79	0.78	1409

```

In [ ]: rf_classifier = RandomForestClassifier(random_state=42)

param_grid = {
    "n_estimators": [50, 100, 200, 300],
    #"max_depth": [3, 5, 7, 9, None],
    #"min_samples_split": [2, 5, 10, 20],
    #"min_samples_leaf": [1, 5, 10, 20],
    #"max_features": ["auto", "sqrt", "log2", None]
}

random_search = RandomizedSearchCV(rf_classifier, param_grid, n_iter=20, cv=5)
random_search.fit(X_train, y_train)

best_model = random_search.best_estimator_

y_pred = best_model.predict(X_test)

report = classification_report(y_test, y_pred)
print(report)

plot_cf_matrix(y_test, y_pred, target_names)

importances = best_model.feature_importances_
indices = np.argsort(importances)[::-1]
feature_names = X.columns

plt.figure(figsize=(12, 6))
plt.title("Feature Importances")
plt.bar(range(X.shape[1]), importances[indices], align="center")

```

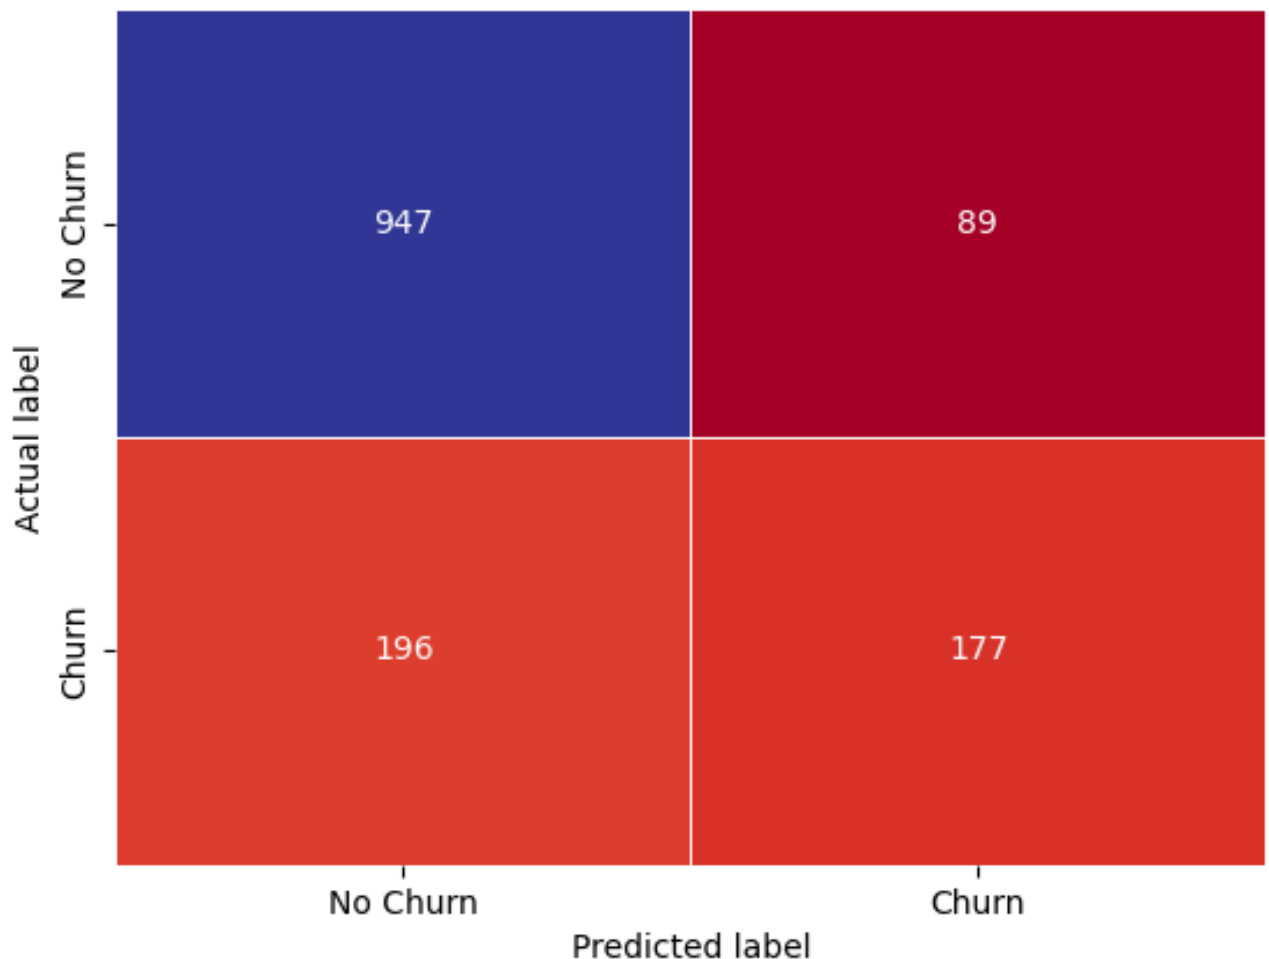
```
plt.xticks(range(X.shape[1]), feature_names[indices], rotation=90)
plt.xlim([-1, X.shape[1]])
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:305: UserWarning: The total space of parameters 4 is smaller than n_iter=20. Running 4 iterations. For exhaustive searches, use GridSearchCV.
```


```
warnings.warn(
```

	precision	recall	f1-score	support
0	0.83	0.91	0.87	1036
1	0.67	0.47	0.55	373
accuracy			0.80	1409
macro avg	0.75	0.69	0.71	1409
weighted avg	0.79	0.80	0.79	1409

Confusion matrix







```
In [ ]: dt_classifier = DecisionTreeClassifier(random_state=42)

param_grid = {
    "max_depth": [3, 5, 7, 9, None],
    "min_samples_split": [2, 5, 10, 20],
    "min_samples_leaf": [1, 5, 10, 20],
    "max_features": ["auto", "sqrt", "log2", None]
}

random_search = RandomizedSearchCV(dt_classifier, param_grid, n_iter=20, cv=5)
random_search.fit(X_train, y_train)

best_model = random_search.best_estimator_

y_pred = best_model.predict(X_test)

report = classification_report(y_test, y_pred)
print(report)

plot_cf_matrix(y_test, y_pred, target_names)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max features='auto'` has been deprecated in 1.1 and will be re
```

```


moved in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.
    warnings.warn(

```

	precision	recall	f1-score	support
0	0.86	0.85	0.86	1036
1	0.61	0.63	0.62	373
accuracy			0.79	1409
macro avg	0.74	0.74	0.74	1409
weighted avg	0.80	0.79	0.79	1409

決策樹混淆矩陣：

Confusion matrix



```
In [ ]: xgb_classifier = XGBClassifier(use_label_encoder=False, eval_metric="logloss")

param_grid = {
    "n_estimators": [50, 100, 200, 300],
    #"max_depth": [3, 5, 7, 9],
    #"learning_rate": [0.01, 0.05, 0.1, 0.2],
    #"subsample": [0.5, 0.75, 1],
    #"colsample_bytree": [0.5, 0.75, 1]
}

random_search = RandomizedSearchCV(xgb_classifier, param_grid, n_iter=20, cv=5)
random_search.fit(X_train, y_train)

best_random_model = random_search.best_estimator_

y_pred = best_random_model.predict(X_test)

report = classification_report(y_test, y_pred)
print("Classification Report:\n", report)
```



Classification Report:					
	precision	recall	f1-score	support	
0	0.84	0.91	0.88	1036	
1	0.68	0.53	0.59	373	
accuracy			0.81	1409	
macro avg	0.76	0.72	0.73	1409	
weighted avg	0.80	0.81	0.80	1409	

Confusion matrix

