

# Coffee Quality Data (CQI May-2023)

## 資料集

In [2]:

```
!gdown --id 1-Snv-zd_aCoFLv2n67hjz42MP77bLHNI
```

```
/usr/local/lib/python3.10/dist-packages/gdown/cli.py:121: FutureWarning: Option `--id` was deprecated in version 4.3.1 and will be removed in 5.0. You don't need to pass it anymore to use a file ID.
```

```
warnings.warn(
```

```
Downloading...
```

```
From: https://drive.google.com/uc?id=1-Snv-zd\_aCoFLv2n67hjz42MP77bLHNI (https://drive.google.com/uc?id=1-Snv-zd\_aCoFLv2n67hjz42MP77bLHNI)
```

```
To: /content/df_arabica_clean.csv
```

```
100% 113k/113k [00:00<00:00, 113MB/s]
```

In [62]:

```
import pandas as pd
```

```
import numpy as np
```

In [4]:


```
col_use = ['Country of Origin', 'Variety', 'Processing Method', 'Color',
           'Aroma', 'Flavor', 'Aftertaste', 'Acidity', 'Body', 'Balance',
           'Uniformity', 'Clean Cup', 'Sweetness', 'Overall', 'Defects',
           'Total Cup Points', 'Moisture Percentage', 'Category One Defects',
           'Quakers', 'Category Two Defects']

df_coffee = pd.read_csv('df_arabica_clean.csv', usecols=col_use)
df_coffee
```

Out[4]:

	Country of Origin	Variety	Processing Method	Aroma	Flavor	Aftertaste	Acidity	Body
0	Colombia	Castillo	Double Anaerobic Washed	8.58	8.50	8.42	8.58	8.25
1	Taiwan	Gesha	Washed / Wet	8.50	8.50	7.92	8.00	7.92
2	Laos	Java	Semi Washed	8.33	8.42	8.08	8.17	7.92
3	Costa Rica	Gesha	Washed / Wet	8.08	8.17	8.17	8.25	8.17
4	Colombia	Red Bourbon	Honey,Mossto	8.33	8.33	8.08	8.25	7.92
...	...	...	...	...	...	...	...	...
202	Brazil	Mundo Novo	Natural / Dry	7.17	7.17	6.92	7.17	7.42
203	Nicaragua	SHG	Natural / Dry	7.33	7.08	6.75	7.17	7.42
204	Laos	Catimor	Washed / Wet	7.25	7.17	7.08	7.00	7.08
205	El Salvador	Maragogype	Natural / Dry	6.50	6.75	6.75	7.17	7.08
206	Brazil	Mundo Novo	SEMI-LAVADO	7.25	7.08	6.67	6.83	6.83

207 rows × 20 columns



變數多時可以先拉第一行出來看一下資料

In [5]:

```
df_coffee.iloc[0]
```

Out[5]:

Country of Origin	Colombia
Variety	Castillo
Processing Method	Double Anaerobic Washed
Aroma	8.58
Flavor	8.5
Aftertaste	8.42
Acidity	8.58
Body	8.25
Balance	8.42
Uniformity	10.0
Clean Cup	10.0
Sweetness	10.0
Overall	8.58
Defects	0.0
Total Cup Points	89.33
Moisture Percentage	11.8
Category One Defects	0
Quakers	0
Color	green
Category Two Defects	3
Name: 0, dtype: object	

In [6]:

```
df_coffee.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 207 entries, 0 to 206
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Country of Origin                    207 non-null   object
1   Variety                              201 non-null   object
2   Processing Method                    202 non-null   object
3   Aroma                               207 non-null   float64
4   Flavor                               207 non-null   float64
5   Aftertaste                           207 non-null   float64
6   Acidity                              207 non-null   float64
7   Body                                 207 non-null   float64
8   Balance                              207 non-null   float64
9   Uniformity                           207 non-null   float64
10  Clean Cup                            207 non-null   float64
11  Sweetness                           207 non-null   float64
12  Overall                              207 non-null   float64
13  Defects                              207 non-null   float64
14  Total Cup Points                     207 non-null   float64
15  Moisture Percentage                  207 non-null   float64
16  Category One Defects                 207 non-null   int64
17  Quakers                             207 non-null   int64
18  Color                                207 non-null   object
19  Category Two Defects                 207 non-null   int64
dtypes: float64(13), int64(3), object(4)
memory usage: 32.5+ KB
```

In [7]:

```
df_coffee.describe()
```

Out[7]:

	Aroma	Flavor	Aftertaste	Acidity	Body	Balance	Unif
count	207.000000	207.000000	207.000000	207.000000	207.000000	207.000000	207.0
mean	7.721063	7.744734	7.599758	7.69029	7.640918	7.644058	9.9
std	0.287626	0.279613	0.275911	0.25951	0.233499	0.256299	0.1
min	6.500000	6.750000	6.670000	6.83000	6.830000	6.670000	8.6
25%	7.580000	7.580000	7.420000	7.50000	7.500000	7.500000	10.0
50%	7.670000	7.750000	7.580000	7.67000	7.670000	7.670000	10.0
75%	7.920000	7.920000	7.750000	7.87500	7.750000	7.790000	10.0
max	8.580000	8.500000	8.420000	8.58000	8.250000	8.420000	10.0

## 清整資料

由上面的基本統計值可以知道 **Defects** 的值都是 0，而 **Clean Cup** 和 **Sweetness** 全部值都是 10。所以這邊先移除這幾個欄位

In [8]:

```
df_coffee.drop(['Defects', 'Clean Cup', 'Sweetness'], axis=1, inplace=True)
```

其中只有品種(Variety) 的部份有一些空值。這邊直接用數量最多的補值

In [9]:

```
from sklearn.impute import SimpleImputer

cat_imputer = SimpleImputer(strategy="most_frequent")

cat_list = df_coffee.select_dtypes(include=["object"]).columns.tolist() # 選擇類別
print('Categorical features:', cat_list)
df_coffee[cat_list] = cat_imputer.fit_transform(df_coffee[cat_list])
```

```
Categorical features: ['Country of Origin', 'Variety', 'Processing Method', 'Color']
```

In [10]:

```
df_coffee.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 207 entries, 0 to 206
```

```
Data columns (total 17 columns):
```

#	Column	Non-Null Count	Dtype
0	Country of Origin	207 non-null	object
1	Variety	207 non-null	object
2	Processing Method	207 non-null	object
3	Aroma	207 non-null	float64
4	Flavor	207 non-null	float64
5	Aftertaste	207 non-null	float64
6	Acidity	207 non-null	float64
7	Body	207 non-null	float64
8	Balance	207 non-null	float64
9	Uniformity	207 non-null	float64
10	Overall	207 non-null	float64
11	Total Cup Points	207 non-null	float64
12	Moisture Percentage	207 non-null	float64
13	Category One Defects	207 non-null	int64
14	Quakers	207 non-null	int64
15	Color	207 non-null	object
16	Category Two Defects	207 non-null	int64

```
dtypes: float64(10), int64(3), object(4)
```

```
memory usage: 27.6+ KB
```

## 視覺化觀察資料

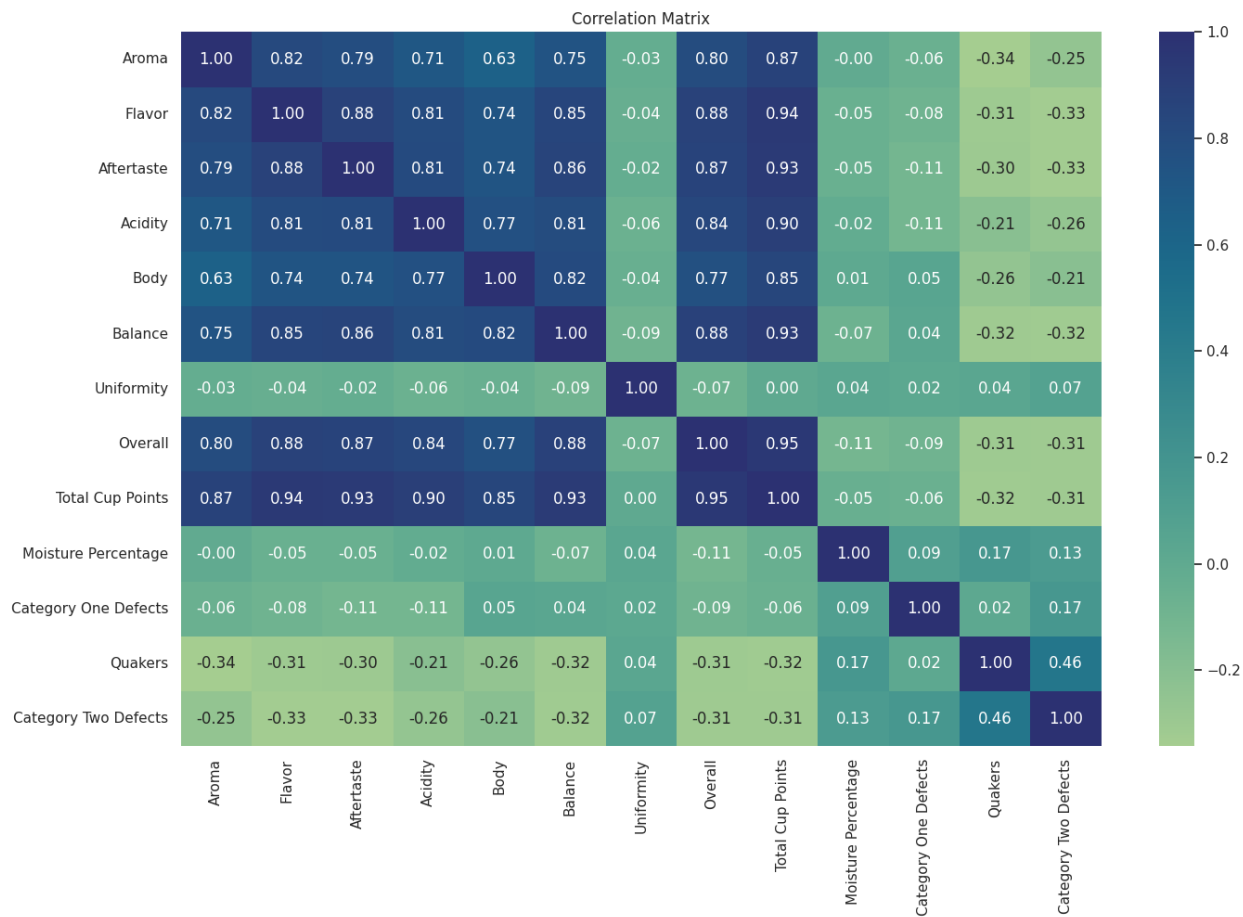
In [11]:

```
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
```

In [12]:

```
corr_matrix = df_coffee.corr()

plt.figure(figsize=(16,10))
sns.heatmap(corr_matrix, annot=True, cbar=True, fmt=".2f", cmap="crest")
plt.title("Correlation Matrix")
plt.show()
```




In [13]:

```
cat_list = df_coffee.select_dtypes(include=["object"]).columns.tolist()

for col in cat_list:
    plt.figure(figsize=(6,4))
    top10 = df_coffee[col].value_counts()[:10]
    top10.plot(kind='bar')
    plt.title("Top 10 " + col)
    plt.grid(visible=False)
    plt.show()
```

Top 10 Country of Origin



## 定義目標與切分 train、test

In [14]:

```
from sklearn.model_selection import train_test_split
```



In [15]:

```
df_coffee["Country of Origin"].value_counts()
```

Out[15]:

Taiwan	61
Guatemala	21
Colombia	19
Honduras	13
Thailand	12
Ethiopia	11
Brazil	10
Costa Rica	8
Nicaragua	7
El Salvador	7
Tanzania, United Republic Of	6
United States (Hawaii)	5
Mexico	4
Peru	4
Vietnam	4
Uganda	3
Indonesia	3
Laos	3
Panama	2
Kenya	2
Madagascar	1
Myanmar	1

Name: Country of Origin, dtype: int64

有些國家太少，這邊取 top 7 (數量 > 10)

In [16]:

```
others_country = df_coffee["Country of Origin"].value_counts()[7:].index
print('other:', others_country)
```

```
other: Index(['Costa Rica', 'Nicaragua', 'El Salvador',
             'Tanzania, United Republic Of', 'United States (Hawaii)', 'M
exico',
             'Peru', 'Vietnam', 'Uganda', 'Indonesia', 'Laos', 'Panama',
             'Kenya',
             'Madagascar', 'Myanmar'],
            dtype='object')
```

In [17]:

```
df_coffee.loc[df_coffee["Country of Origin"].isin(others_country), "Country of Or:
```

看一下處理完的 y 的數量

In [18]:

```
df_coffee["Country of Origin"].value_counts()
```

Out[18]:

```
Taiwan      61
others      60
Guatemala   21
Colombia     19
Honduras     13
Thailand     12
Ethiopia     11
Brazil       10
Name: Country of Origin, dtype: int64
```

切分資料

In [112]:

```
X = df_coffee.drop(labels="Country of Origin", axis=1)
y = df_coffee["Country of Origin"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

這邊要注意的是，sklearn 的 `train_test_split` 會根據你丟進去的 `data type` 去輸出對應格式的切分結果。如果我 `X`、`y` 給他 `array` 的格式，最後輸出的也是 `array`。在這邊我們直接使用 `pandas dataframe` 格式，他會輸出 `dataframe` 的格式，我就可以透過 `index` 去追朔切分完的資料是對應哪筆原始的資料。

In [20]:

```
X_train
```

Out[20]:

	Variety	Processing Method	Aroma	Flavor	Aftertaste	Acidity	Body	Balance	Unif
34	Ethiopian Heirlooms	Natural / Dry	8.00	8.08	8.00	8.00	7.67	7.75	
132	Lempira	Washed / Wet	7.67	7.75	7.67	7.58	7.50	7.58	
123	Typica	Washed / Wet	7.58	7.67	7.58	7.75	7.42	7.58	
156	Bourbon	Washed / Wet	7.67	7.33	7.17	7.67	7.83	7.50	
175	Catimor	Washed / Wet	7.17	7.50	7.42	7.42	7.58	7.50	
...	...	...	...	...	...	...	...	...	
140	Typica	Pulped natural / honey	7.50	7.67	7.58	7.58	7.50	7.50	
102	Typica	Washed / Wet	7.67	7.92	7.58	7.58	7.58	7.75	
64	Catrenic	Natural / Dry	7.92	7.75	7.67	7.67	7.83	7.75	
56	Yellow Bourbon	Double Carbonic Maceration / Natural	7.83	7.92	7.75	7.92	7.67	7.83	
190	Java	Washed / Wet	7.33	7.42	7.25	7.33	7.50	7.42	

165 rows × 16 columns



可以看到這樣就可以，透過 index 直接找到對應的資料。

In [21]:

```
df_coffee.iloc[X_train.index]
```

Out[21]:

	Country of Origin	Variety	Processing Method	Aroma	Flavor	Aftertaste	Acidity	Body	Balance
34	Ethiopia	Ethiopian Heirlooms	Natural / Dry	8.00	8.08	8.00	8.00	7.67	
132	Honduras	Lempira	Washed / Wet	7.67	7.75	7.67	7.58	7.50	
123	Taiwan	Typica	Washed / Wet	7.58	7.67	7.58	7.75	7.42	
156	others	Bourbon	Washed / Wet	7.67	7.33	7.17	7.67	7.83	
175	Thailand	Catimor	Washed / Wet	7.17	7.50	7.42	7.42	7.58	
...	...	...	...	...	...	...	...	...	
140	Taiwan	Typica	Pulped natural / honey	7.50	7.67	7.58	7.58	7.50	
102	others	Typica	Washed / Wet	7.67	7.92	7.58	7.58	7.58	
64	others	Catrenic	Natural / Dry	7.92	7.75	7.67	7.67	7.83	
56	Brazil	Yellow Bourbon	Double Carbonic Maceration / Natural	7.83	7.92	7.75	7.92	7.67	
190	Thailand	Java	Washed / Wet	7.33	7.42	7.25	7.33	7.50	

165 rows × 17 columns



## 資料轉換

ColumnTransformer 可以結合 sklearn pipeline 的特性，對不同 columns 建立各自的 pipeline

[sklearn.compose.ColumnTransformer — scikit-learn 1.2.2 documentation \(https://scikit-learn.org/stable/modules/generated/sklearn.compose.ColumnTransformer.html\)](https://scikit-learn.org/stable/modules/generated/sklearn.compose.ColumnTransformer.html)

In [22]:

```
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, MinMaxScaler
```

In [56]:

```
X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 165 entries, 34 to 190
Data columns (total 16 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Variety                              165 non-null    object
 1   Processing Method                    165 non-null    object
 2   Aroma                                165 non-null    float64
 3   Flavor                               165 non-null    float64
 4   Aftertaste                           165 non-null    float64
 5   Acidity                              165 non-null    float64
 6   Body                                  165 non-null    float64
 7   Balance                              165 non-null    float64
 8   Uniformity                           165 non-null    float64
 9   Overall                              165 non-null    float64
10   Total Cup Points                     165 non-null    float64
11   Moisture Percentage                  165 non-null    float64
12   Category One Defects                 165 non-null    int64
13   Quakers                             165 non-null    int64
14   Color                                165 non-null    object
15   Category Two Defects                 165 non-null    int64
dtypes: float64(10), int64(3), object(3)
memory usage: 21.9+ KB
```

In [24]:

```
cat_list = X_train.select_dtypes(include=["object"]).columns.tolist()
num_list = X_train.select_dtypes(exclude=["object"]).columns.tolist()

full_pipeline = ColumnTransformer([
    ("num", MinMaxScaler(), num_list),
    ("category", OneHotEncoder(handle_unknown="ignore"), cat_list),
])
```

In [25]:

```
X_train_norm = full_pipeline.fit_transform(X_train)
X_test_norm = full_pipeline.transform(X_test)
X_norm = full_pipeline.transform(X)
```

自動轉換為稀疏矩陣的資料格式

In [26]:

```
X_train_norm
```

Out[26]:

```
<165x73 sparse matrix of type '<class 'numpy.float64'>'
  with 2290 stored elements in Compressed Sparse Row format>
```

In [27]:

```
print(X_train_norm)
```

(0, 0)	0.7211538461538463
(0, 1)	0.7599999999999998
(0, 2)	0.7599999999999998
(0, 3)	0.6685714285714286
(0, 4)	0.591549295774648
(0, 5)	0.6171428571428565
(0, 6)	1.0
(0, 7)	0.6073298429319371
(0, 8)	0.6469549867608118
(0, 9)	0.9111111111111111
(0, 11)	0.25
(0, 12)	0.3076923076923077
(0, 28)	1.0
(0, 57)	1.0
(0, 70)	1.0
(1, 0)	0.5625
(1, 1)	0.5714285714285712
(1, 2)	0.5714285714285712
(1, 3)	0.4285714285714284
(1, 4)	0.471830985915493
(1, 5)	0.5199999999999996
(1, 6)	1.0
(1, 7)	0.4345549738219896
(1, 8)	0.4633715798764344
(1, 9)	0.7777777777777777
:	:
(163, 5)	0.6628571428571428
(163, 6)	1.0
(163, 7)	0.6073298429319371
(163, 8)	0.5957634598411294
(163, 9)	0.8296296296296295
(163, 11)	0.08333333333333333
(163, 12)	0.38461538461538464
(163, 50)	1.0
(163, 55)	1.0
(163, 66)	1.0
(164, 0)	0.3990384615384617
(164, 1)	0.38285714285714256
(164, 2)	0.33142857142857096
(164, 3)	0.28571428571428603
(164, 4)	0.471830985915493
(164, 5)	0.42857142857142794
(164, 6)	1.0
(164, 7)	0.26178010471204205
(164, 8)	0.3018534863195059
(164, 9)	0.8592592592592592
(164, 10)	0.5
(164, 12)	0.07692307692307693
(164, 31)	1.0
(164, 61)	1.0
(164, 66)	1.0



# 模型預測: random forest

In [29]:

```
from sklearn.ensemble import RandomForestClassifier
```

```
clf = RandomForestClassifier(random_state=0)
```

```
clf.fit(X_train_norm, y_train)
```

```
rf_pred = clf.predict(X_test_norm)
```

```
print(rf_pred)
```

```
['Colombia' 'others' 'others' 'others' 'Colombia' 'others' 'Guatemala'
la'
'others' 'Taiwan' 'Ethiopia' 'others' 'Honduras' 'Taiwan' 'Taiwan'
'others' 'others' 'Taiwan' 'Taiwan' 'others' 'Taiwan' 'Taiwan'
'Guatemala' 'Taiwan' 'Taiwan' 'Colombia' 'Taiwan' 'others' 'Taiwan'
n'
'Guatemala' 'others' 'others' 'Taiwan' 'others' 'Taiwan' 'Taiwan'
'Ethiopia' 'others' 'others' 'Taiwan' 'Taiwan' 'Taiwan' 'others']
```

In [30]:

```
from sklearn.metrics import classification_report
```

```
print(classification_report(y_test, rf_pred))
```

	precision	recall	f1-score	support
Brazil	0.00	0.00	0.00	1
Colombia	0.67	0.33	0.44	6
Ethiopia	1.00	0.50	0.67	4
Guatemala	0.33	0.17	0.22	6
Honduras	1.00	0.50	0.67	2
Taiwan	0.65	1.00	0.79	11
Thailand	0.00	0.00	0.00	2
others	0.50	0.80	0.62	10
accuracy			0.60	42
macro avg	0.52	0.41	0.43	42
weighted avg	0.57	0.60	0.54	42

# 模型預測: kmeans (k=目標類別數)

In [31]:

```
from sklearn.cluster import KMeans
kmeans = KMeans(
    init="random",
    n_clusters=7,
    n_init='auto',
    random_state=42
)
kmeans.fit(X_norm)
```

Out[31]:

KMeans(init='random', n\_clusters=7, n\_init='auto', random\_state=42)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

要注意的是這邊出來的是全部的結果(train + test) · 需要把對應 test 結果抽取出來才能做比較。這時候就需要使用到上面切分 train、test 時對應原始資料的 index

In [32]:

```
kmeans.labels_
```

Out[32]:

```
array([5, 5, 5, 5, 5, 5, 5, 6, 0, 0, 6, 5, 5, 5, 2, 4, 1, 5, 5, 0,
5, 6,
      4, 3, 3, 2, 5, 5, 5, 1, 3, 1, 5, 1, 6, 0, 3, 5, 0, 3, 6, 0,
5, 1,
      1, 5, 2, 1, 2, 1, 6, 3, 0, 0, 6, 1, 1, 3, 4, 2, 6, 2, 4, 5,
6, 5,
      6, 6, 5, 4, 1, 1, 6, 3, 1, 3, 6, 4, 1, 4, 1, 3, 1, 2, 4, 5,
2, 2,
      4, 0, 1, 0, 2, 2, 5, 1, 6, 2, 6, 0, 1, 2, 1, 4, 1, 0, 3, 6,
4, 6,
      6, 4, 2, 1, 3, 1, 0, 1, 1, 3, 1, 4, 1, 1, 0, 2, 3, 3, 0, 6,
1, 4,
      3, 1, 6, 3, 6, 1, 1, 3, 4, 1, 0, 3, 1, 3, 4, 6, 4, 2, 0, 6,
1, 6,
      2, 1, 1, 2, 6, 3, 3, 1, 1, 3, 3, 1, 3, 6, 0, 1, 1, 1, 3, 4,
1, 1,
      2, 2, 4, 0, 0, 4, 4, 4, 6, 3, 1, 2, 4, 4, 1, 6, 1, 0, 6, 6,
0, 1,
      0, 0, 6, 6, 6, 6, 1, 6, 1], dtype=int32)
```

In [44]:

```
kmeans_train_pred = kmeans.labels_[X_train.index]
```

需要先依照 train 的結果將 kmean 預測的 group 轉換為對應類別

In [46]:

```
map_df = pd.DataFrame({'kmeans': kmeans_train_pred, 'label': y_train})
map_df
```

Out[46]:

	kmeans	label
34	6	Ethiopia
132	3	Honduras
123	1	Taiwan
156	1	others
175	1	Thailand
...	...	...
140	4	Taiwan
102	1	others
64	6	others
56	1	Brazil
190	1	Thailand

165 rows × 2 columns

In [55]:

```
map_dt={}
for gp in map_df['kmeans'].unique():
    most_freq_label = map_df.loc[map_df['kmeans'] == gp, 'label'].value_counts().idxmax()
    map_dt.update({gp: most_freq_label})
```

可以看到各 group 對應的標籤，有些標籤可能沒有，這是正常的。

In [54]:

```
map_dt
```

Out[54]:

```
{6: 'Taiwan',  
 3: 'others',  
 1: 'others',  
 2: 'Guatemala',  
 0: 'others',  
 4: 'Taiwan',  
 5: 'Taiwan'}
```

In [66]:

```
kmeans_test_pred = kmeans.labels_[X_test.index]  
kmeans_test_pred
```

Out[66]:

```
array([2, 2, 1, 6, 1, 6, 1, 6, 3, 6, 2, 1, 4, 1, 4, 4, 3, 6, 6, 0,  
       5, 6,  
       1, 4, 0, 0, 3, 5, 5, 1, 4, 4, 4, 5, 5, 0, 2, 0, 5, 5, 3, 0],  
      dtype=int32)
```

In [67]:

```
kmeans_test_pred = np.array([map_dt[i] for i in kmeans_test_pred])  
kmeans_test_pred
```

Out[67]:

```
array(['Guatemala', 'Guatemala', 'others', 'Taiwan', 'others', 'Tai  
wan',  
      'others', 'Taiwan', 'others', 'Taiwan', 'Guatemala', 'other  
s',  
      'Taiwan', 'others', 'Taiwan', 'Taiwan', 'others', 'Taiwan',  
      'Taiwan', 'others', 'Taiwan', 'Taiwan', 'others', 'Taiwan',  
      'others', 'others', 'others', 'Taiwan', 'Taiwan', 'others',  
      'Taiwan', 'Taiwan', 'Taiwan', 'Taiwan', 'Taiwan', 'others',  
      'Guatemala', 'others', 'Taiwan', 'Taiwan', 'others', 'other  
s'],  
      dtype='<U9')
```

轉換完成後就可以評估我們 kmeans 的結果

In [68]:

```
from sklearn.metrics import classification_report

print(classification_report(y_test, kmeans_test_pred))
```

	precision	recall	f1-score	support
Brazil	0.00	0.00	0.00	1
Colombia	0.00	0.00	0.00	6
Ethiopia	0.00	0.00	0.00	4
Guatemala	0.25	0.17	0.20	6
Honduras	0.00	0.00	0.00	2
Taiwan	0.33	0.64	0.44	11
Thailand	0.00	0.00	0.00	2
others	0.24	0.40	0.30	10
accuracy			0.29	42
macro avg	0.10	0.15	0.12	42
weighted avg	0.18	0.29	0.21	42

## 模型預測: kmeans (elbow 選擇 k)

In [75]:


```
from sklearn.cluster import KMeans

kmeans_kwargs = {
    "init": "random",
    "random_state": 42,
}

# A list holds the SSE values for each k
sse = []
for k in range(1, 11):
    kmeans = KMeans(
        n_clusters=k,
        init="random",
        n_init='auto',
        random_state=42
    )
    kmeans.fit(X_norm)
    sse.append(kmeans.inertia_)
```

In [76]:

```
plt.style.use("fivethirtyeight")
plt.plot(range(1, 11), sse)
plt.xticks(range(1, 11))
plt.xlabel("Number of Clusters")
plt.ylabel("SSE")
plt.show()
```



In [72]:

```
!pip install kneed
```

Looking in indexes: <https://pypi.org/simple>, (<https://pypi.org/simple>,) <https://us-python.pkg.dev/colab-wheels/public/simple/> (<https://us-python.pkg.dev/colab-wheels/public/simple/>)

Collecting kneed

Downloading kneed-0.8.3-py3-none-any.whl (10 kB)

Requirement already satisfied: numpy>=1.14.2 in /usr/local/lib/python3.10/dist-packages (from kneed) (1.22.4)

Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from kneed) (1.10.1)

Installing collected packages: kneed

Successfully installed kneed-0.8.3

In [77]:

```
from kneed import KneeLocator

kl = KneeLocator(
    range(1, 11), sse, curve="convex", direction="decreasing"
)

kl.elbow
```

Out[77]:

4

使用 k = 4

In [78]:

```
kmeans = KMeans(
    n_clusters=4,
    init="random",
    n_init='auto',
    random_state=42
)
kmeans.fit(X_norm)
```

Out[78]:

KMeans(init='random', n\_clusters=4, n\_init='auto', random\_state=42)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [79]:

```
kmeans_train_pred = kmeans.labels_[X_train.index]
map_df = pd.DataFrame({'kmeans': kmeans_train_pred, 'label': y_train})

map_dt={}
for gp in map_df['kmeans'].unique():
    most_freq_label = map_df.loc[map_df['kmeans'] == gp, 'label'].value_counts().max()
    map_dt.update({gp: most_freq_label})

kmeans_test_pred = kmeans.labels_[X_test.index]
kmeans_test_pred = np.array([map_dt[i] for i in kmeans_test_pred])
kmeans_test_pred
```

Out[79]:

```
array(['others', 'others', 'others', 'Taiwan', 'others', 'Taiwan',
      'others', 'Taiwan', 'others', 'Taiwan', 'others', 'others',
      'Taiwan', 'others', 'Taiwan', 'Taiwan', 'others', 'Taiwan',
      'Taiwan', 'others', 'others', 'Taiwan', 'others', 'Taiwan',
      'others', 'others', 'others', 'Taiwan', 'Taiwan', 'others',
      'Taiwan', 'Taiwan', 'Taiwan', 'Taiwan', 'others', 'others',
      'others', 'others', 'Taiwan', 'others', 'others', 'others'],
      dtype='<U6')
```

In [80]:

```
from sklearn.metrics import classification_report

print(classification_report(y_test, kmeans_test_pred))
```

	precision	recall	f1-score	support
Brazil	0.00	0.00	0.00	1
Colombia	0.00	0.00	0.00	6
Ethiopia	0.00	0.00	0.00	4
Guatemala	0.00	0.00	0.00	6
Honduras	0.00	0.00	0.00	2
Taiwan	0.22	0.36	0.28	11
Thailand	0.00	0.00	0.00	2
others	0.21	0.50	0.29	10
accuracy			0.21	42
macro avg	0.05	0.11	0.07	42
weighted avg	0.11	0.21	0.14	42

## 模型預測: random forest + kmeans



In [81]:

```
kmeans = KMeans(  
    n_clusters=4,  
    init="random",  
    n_init='auto',  
    random_state=42  
)  
kmeans.fit(X_norm)
```

Out[81]:

```
KMeans(init='random', n_clusters=4, n_init='auto', random_state=42)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with [nbviewer.org](https://nbviewer.org).**

In [93]:

```
kmeans_train_pred = kmeans.labels_[X_train.index]  
X_train['kmean_label'] = kmeans_train_pred
```

In [94]:

```
kmeans_test_pred = kmeans.labels_[X_test.index]  
X_test['kmean_label'] = kmeans_test_pred
```

這邊直接用上面 資料轉換 地方定義的 full pipeline 一樣的方式進行正規化、onehot encoding

In [89]:

```
cat_list.append('kmean_label')  
cat_list
```

Out[89]:

```
['Variety', 'Processing Method', 'Color', 'kmean_label']
```

In [95]:

```
full_pipeline_kmean = ColumnTransformer([  
    ("num", MinMaxScaler(), num_list),  
    ("category", OneHotEncoder(handle_unknown="ignore"), cat_list),  
)
```

In [96]:

```
X_train_kmean_norm = full_pipeline_kmean.fit_transform(X_train)
X_test_kmean_norm = full_pipeline_kmean.transform(X_test)
```

In [99]:

```
from sklearn.ensemble import RandomForestClassifier

clf = RandomForestClassifier(random_state=0)
clf.fit(X_train_kmean_norm, y_train)

rf_kmean_pred = clf.predict(X_test_kmean_norm)
```

In [100]:

```
from sklearn.metrics import classification_report

print(classification_report(y_test, rf_kmean_pred))
```

	precision	recall	f1-score	support
Brazil	0.00	0.00	0.00	1
Colombia	0.67	0.33	0.44	6
Ethiopia	1.00	0.50	0.67	4
Guatemala	0.33	0.17	0.22	6
Honduras	1.00	0.50	0.67	2
Taiwan	0.67	0.91	0.77	11
Thailand	0.00	0.00	0.00	2
others	0.50	0.90	0.64	10
accuracy			0.60	42
macro avg	0.52	0.41	0.43	42
weighted avg	0.58	0.60	0.55	42